

✓ CSET419 – Introduction to Generative AI

Lab – 5

Baseline CNN for Image-to-Image Translation (Encoder–Decoder without GAN)

🎯 Objective

To implement a baseline Encoder–Decoder Convolutional Neural Network (CNN) for paired image-to-image translation and analyze its performance using reconstruction loss.

📌 Experiment Tasks

1. Load paired images (CIFAR10 dataset)
 2. Normalize images to range [-1, 1]
 3. Train an Encoder–Decoder CNN
 4. Compute reconstruction loss (MSE / L1)
 5. Visualize translated (reconstructed) images
-

📁 Dataset

CIFAR10 Dataset

- 60,000 RGB images
 - Image size: 32×32
 - 10 object classes
 - Used in paired setting (Input = Target image)
-

🧠 Model Architecture

Encoder

- Conv2D ($3 \rightarrow 64$), ReLU
- Conv2D ($64 \rightarrow 128$), ReLU

Decoder

- ConvTranspose2D ($128 \rightarrow 64$), ReLU
- ConvTranspose2D ($64 \rightarrow 3$), Tanh

The final activation is **Tanh** to ensure output is in range [-1, 1].

📉 Loss Function

Reconstruction Loss:

- Mean Squared Error (MSE)
- or
- L1 Loss

The model minimizes the difference between the input image and reconstructed image.

⚙️ Training Configuration

- Optimizer: Adam
 - Learning Rate: 0.001
 - Epochs: 5
 - Batch Size: 16
 - Device: CPU / GPU
-

📊 Results & Observation

- The reconstructed images appear slightly blurry.
- Pixel-wise reconstruction losses (MSE/L1) cause smoothing.

- No adversarial training is used.

✓ Conclusion

The baseline Encoder–Decoder CNN successfully reconstructs CIFAR10 images.

However, outputs are blurry due to pixel-wise loss limitations.

To improve results:

- Use U-Net (skip connections)
 - Use GAN-based training
 - Use perceptual loss
- ```
=====
CSET419 - Introduction to Generative AI
Lab - 5
Baseline CNN for Image-to-Image Translation
Encoder-Decoder without GAN
=====

import torch
import torch.nn as nn
import torch.optim as optim
import torchvision
import torchvision.transforms as transforms
import matplotlib.pyplot as plt

1. Device Configuration

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print("Using device:", device)

2. Load CIFAR10 Dataset (Paired Setting)
Normalize images to [-1, 1]

transform = transforms.Compose([
 transforms.ToTensor(),
 transforms.Normalize((0.5, 0.5, 0.5),
 (0.5, 0.5, 0.5)) # Scale to [-1,1]
])

dataset = torchvision.datasets.CIFAR10(
 root='./data',
 train=True,
 download=True,
 transform=transform
)

dataloader = torch.utils.data.DataLoader(
 dataset,
 batch_size=16,
 shuffle=True
)

3. Define Encoder-Decoder CNN

class EncoderDecoder(nn.Module):
 def __init__(self):
 super(EncoderDecoder, self).__init__()

 # Encoder
 self.encoder = nn.Sequential(
 nn.Conv2d(3, 64, kernel_size=4, stride=2, padding=1),
 nn.ReLU(True),

 nn.Conv2d(64, 128, kernel_size=4, stride=2, padding=1),
 nn.ReLU(True)
)

 # Decoder
 self.decoder = nn.Sequential(
 nn.ConvTranspose2d(128, 64, kernel_size=4, stride=2, padding=1),
 nn.ReLU(True),

 nn.ConvTranspose2d(64, 3, kernel_size=4, stride=2, padding=1),
 nn.Tanh() # Output in [-1,1]
)
```

```
def forward(self, x):
 x = self.encoder(x)
 x = self.decoder(x)
 return x

model = EncoderDecoder().to(device)

4. Loss and Optimizer

criterion = nn.MSELoss() # Reconstruction loss
optimizer = optim.Adam(model.parameters(), lr=0.001)

5. Training Loop

epochs = 5

for epoch in range(epochs):
 running_loss = 0.0

 for images, _ in dataloader:
 images = images.to(device)

 optimizer.zero_grad()

 outputs = model(images)

 # Paired translation: Input = Target
 loss = criterion(outputs, images)

 loss.backward()
 optimizer.step()

 running_loss += loss.item()

 print(f"Epoch [{epoch+1}/{epochs}] Loss: {running_loss/len(dataloader):.4f}")

print("Training Completed")

6. Clear Visualization
Input | Expected | Output

model.eval()

with torch.no_grad():
 sample_images, _ = next(iter(dataloader))
 sample_images = sample_images.to(device)
 outputs = model(sample_images)

 # Convert from [-1,1] back to [0,1]
 inputs = (sample_images + 1) / 2
 expected = inputs.clone() # Ground truth = input (reconstruction task)
 outputs = (outputs + 1) / 2

 inputs = inputs.cpu()
 expected = expected.cpu()
 outputs = outputs.cpu()

 num_images = 5

 plt.figure(figsize=(15, 8))

 for i in range(num_images):

 # ----- Input -----
 plt.subplot(3, num_images, i + 1)
 plt.imshow(inputs[i].permute(1, 2, 0))
 plt.title("Input")
 plt.axis("off")

 # ----- Expected -----
 plt.subplot(3, num_images, num_images + i + 1)
 plt.imshow(expected[i].permute(1, 2, 0))
 plt.title("Expected")
 plt.axis("off")

 # ----- Output -----
 plt.subplot(3, num_images, 2*num_images + i + 1)
 plt.imshow(outputs[i].permute(1, 2, 0))
 plt.title("Output")
```

```
plt.axis("off")

plt.tight_layout()
plt.show()

======
Explanation:
Input = Real CIFAR10 Image
Expected = Same Image (Reconstruction Task)
Output = Model Reconstructed Image (Slightly Blurry)
======
```

```
Using device: cuda
Epoch [1/5] Loss: 0.0038
Epoch [2/5] Loss: 0.0009
Epoch [3/5] Loss: 0.0006
Epoch [4/5] Loss: 0.0005
Epoch [5/5] Loss: 0.0004
Training Completed
```



